

## MACHINE LEARNING

**1) R-squared or Residual Sum of Squares (RSS) which one of these two is a better measure of goodness of fit model in regression and why?**

R-squared is generally considered a better measure of the goodness of fit for a regression model compared to Residual Sum of Squares (RSS). Here's why:

1. **Interpretation:** R-squared measures the proportion of the variance in the dependent variable that is predictable from the independent variables. Higher R-squared values indicate that the regression model fits the data well.
2. **Normalization:** R-squared is normalized, ranging from 0 to 1, which makes it easier to interpret across different datasets and models. It provides a standardized measure of how well the model fits the data.
3. **Comparison:** While RSS measures the total sum of squared residuals (errors) between the predicted and observed values, it doesn't account for the total variation in the dependent variable. Therefore, RSS alone doesn't give a complete picture of the model's goodness of fit.
4. **Adjusted R-squared:** In cases where the number of predictors increases, Adjusted R-squared is preferred over R-squared as it penalizes excessive complexity, providing a more accurate assessment of model fit.

In summary, R-squared is preferred because it directly assesses how well the model explains the variation in the data and provides a standardized measure for comparison across different models and datasets.

**2) What are TSS (Total Sum of Squares), ESS (Explained Sum of Squares) and RSS (Residual Sum of Squares) in regression. Also mention the equation relating these three metrics with each other**

In regression analysis, TSS (Total Sum of Squares), ESS (Explained Sum of Squares), and RSS (Residual Sum of Squares) are important metrics that help evaluate the goodness of fit of a regression model:

1. **Total Sum of Squares (TSS):** TSS represents the total variation in the dependent variable (Y) and is calculated as the sum of the squared differences between each observed dependent variable value and the mean of the dependent variable:

$$TSS = \sum (Y_1 - Y_2)^2$$

where  $Y_1$  are the observed values of the dependent variable and  $Y_2$  is the mean of the dependent variable.

2. **Explained Sum of Squares (ESS):** ESS measures the variation in the dependent variable that is explained by the regression model. It is calculated as the sum of the squared differences between the predicted values ( $\hat{Y}_3$ ) from the regression model and the mean of the dependent variable:

$$ESS = \sum (\hat{Y}_3 - \bar{Y}_2)^2$$

where  $\hat{Y}_3$  are the predicted values from the regression model and  $\bar{Y}_2$  is the mean of the dependent variable.

3. **Residual Sum of Squares (RSS):** RSS quantifies the variation in the dependent variable that is not explained by the regression model, i.e., the residual variation. It is calculated as the sum of the squared differences between the observed values of the dependent variable and the predicted values from the regression model:

$$RSS = \sum (Y_1 - \hat{Y}_3)^2$$

where  $Y_1$  are the observed values and  $\hat{Y}_3$  are the predicted values.

The relationship between these metrics is given by:  $TSS = ESS + RSS$

This equation illustrates that the total variation in the dependent variable (TSS) is partitioned into the variation explained by the model (ESS) and the unexplained residual variation (RSS).

### 3) What is the need of regularization in machine learning?

Regularization is crucial in machine learning to address the problem of overfitting, where a model learns too much from the training data and fails to generalize well to unseen data. The main reasons for using regularization include:

1. **Preventing Overfitting:** Regularization techniques penalize complex models that fit the training data too closely. By imposing constraints on model parameters, regularization helps in creating simpler models that generalize better.
2. **Improving Generalization:** Regularization encourages models to focus on the most important patterns in the data rather than noise, thereby improving their ability to make accurate predictions on new, unseen data.
3. **Handling Multicollinearity:** In regression models with highly correlated predictors, regularization methods like Ridge Regression can stabilize parameter estimates and mitigate the effects of multicollinearity.
4. **Enhancing Model Interpretability:** By reducing model complexity, regularization can make it easier to interpret the

influence of different features on model predictions, thus improving transparency.

5. **Robustness to Outliers:** Some regularization techniques, such as Lasso Regression, can effectively reduce the impact of outliers in the training data, leading to more robust models.

Regularization plays a vital role across various machine learning algorithms, ensuring that models generalize well and perform reliably on new data.

#### 4) What is Gini-impurity index?

The **Gini-impurity index**, is a measure of impurity or homogeneity used in decision tree algorithms during the construction of the tree. It quantifies the likelihood of incorrectly classifying a randomly chosen element if it were randomly labeled according to the distribution of labels in the node.

It calculates the impurity of a node in a decision tree where each class (or label) is represented proportionally to its occurrence. It ranges from 0 to 1, where 0 indicates that the node is pure (all elements belong to the same class), and 1 indicates maximum impurity (elements are evenly distributed across all classes).

#### 5) Are unregularized decision-trees prone to overfitting? If yes, why?

Yes, unregularized decision trees are prone to overfitting. Here's why:

1. **Complexity and Flexibility:** Decision trees are highly flexible and can capture intricate patterns in the training data. Without regularization, they tend to grow deep and complex, fitting the training data very closely.
2. **Capturing Noise:** They tend to capture noise or irrelevant details in the training data, which are specific to the training set but do not generalize well to new, unseen data.
3. **Overfitting:** This phenomenon occurs when a decision tree model becomes overly complex, fitting not only the underlying patterns but also the random noise in the training data. As a result, the model performs exceptionally well on the training data but fails to generalize to new data points.
4. **No Constraints:** Unregularized decision trees do not have constraints on their growth, such as limiting the depth of the tree or the number of leaf nodes. This lack of constraints allows them to become too specific to the training data, leading to overfitting.

Regularization techniques like pruning or limiting the tree depth are used to mitigate these issues by restraining the growth of decision trees and promoting models that generalize better to unseen data.

## 6) What is an ensemble technique in machine learning?

An ensemble technique in machine learning refers to a method that combines multiple individual models together to obtain a more accurate and robust predictive model than any of the individual models alone. The key idea behind ensemble methods is to leverage the diversity among the models to improve overall prediction performance and generalizability.

Ensemble techniques typically involve training multiple base models on the same data and then combining their predictions in various ways. Common ensemble methods include:

1. **Bagging (Bootstrap Aggregating)**
2. **Boosting**
3. **Stacking:**

## 7) What is the difference between Bagging and Boosting techniques?

1. **Bagging (Bootstrap Aggregating):** It involves training multiple instances of the same learning algorithm on different subsets of the training data and averaging the predictions.
2. **Boosting:** This method sequentially builds a series of weak learners, where each subsequent learner corrects errors made by its predecessor, thereby improving the overall predictive performance.

**Bagging** focuses on reducing variance by averaging independent models, whereas **Boosting** focuses on reducing bias by sequentially improving the predictive capability of the models.

## 8) What is out-of-bag error in random forests?

Out-of-bag (OOB) error in random forests is an estimation of how well a random forest model will generalize to unseen data. OOB error is computationally efficient and leverages all available data, making it particularly useful for datasets with limited samples.

## 9) What is K-fold cross-validation?

K-fold cross-validation is a technique used in machine learning to assess the performance of a model. Here's how it works:

1. **Partitioning the Data:** The dataset is divided into K subsets (folds) of equal size. Each subset is used as a validation set once while the remaining K-1 subsets are used as training data.
2. **Iterative Training and Validation:** The model is trained K times, each time using a different fold as the validation set and the remaining folds as training data.

3. **Performance Evaluation:** After training the model K times, performance metrics such as accuracy, precision, or error are computed by averaging the results from all K folds. This provides a more reliable estimate of model performance compared to a single train-test split.
4. **Advantages:** K-fold cross-validation ensures that every data point is used for both training and validation, reducing bias that may result from a single train-test split. It also helps in assessing the model's ability to generalize to new data.
5. **Implementation:** Common values for K are 5 and 10, but it can vary depending on the dataset size and computational resources.

#### 10) What is hyper parameter tuning in machine learning and why it is done?

Hyperparameter tuning in machine learning refers to the process of optimizing the hyperparameters of a model to improve its performance. Hyperparameters are parameters that are set before the learning process begins and cannot be directly learned from the data. Examples include the learning rate in neural networks, the depth of decision trees, or the number of clusters in clustering algorithms.

Hyperparameter tuning helps in:

1. **Optimizing Model Performance:** Hyperparameters significantly impact the model's ability to learn from the data and generalize to new data. Tuning them can lead to better accuracy, precision, or other performance metrics.
2. **Avoiding Overfitting or Underfitting:** By finding optimal hyperparameters, you can prevent the model from either memorizing the training data (overfitting) or failing to capture its patterns (underfitting).
3. **Enhancing Efficiency:** Tuned hyperparameters can make the model training process faster and more efficient, reducing computational costs and time.
4. **Improving Generalization:** Models with well-tuned hyperparameters are more likely to generalize well to unseen data, making them more robust and reliable in real-world applications.
5. **Ensuring Model Stability:** Hyperparameter tuning helps in stabilizing the model's performance across different datasets or different runs, ensuring consistent results.

#### 11) What issues can occur if we have a large learning rate in Gradient Descent?

Having a large learning rate in Gradient Descent can lead to several issues:

1. **Divergence:** When the learning rate is too large, the updates to the model parameters can become too drastic. This can cause the optimization process to overshoot the optimal point and lead to divergence instead of convergence.
2. **Oscillations:** Large learning rates can cause the gradient descent algorithm to oscillate around the minimum rather than converging smoothly. This oscillation can make it difficult for the algorithm to settle on the optimal values of the parameters.
3. **Slow Convergence:** Paradoxically, a learning rate that is too large can slow down the convergence of the algorithm. This happens because the steps taken are so large that the algorithm may keep overshooting the minimum, requiring more iterations to converge.
4. **Difficulty in Finding Optimal Solutions:** A large learning rate can make it harder for the algorithm to find the global minimum or even a good local minimum. It may get stuck in a suboptimal solution due to the overly aggressive updates.

## 12) Can we use Logistic Regression for classification of Non-Linear Data? If not, why?

No, Logistic Regression is not suitable for classifying non-linear data. Because:

1. **Linear Decision Boundary:** Logistic Regression models assume a linear relationship between the input variables and the log-odds of the output. This implies that it can only learn and represent linear decision boundaries between classes [3].
2. **Limited Complexity:** Since Logistic Regression uses a sigmoid function to map predicted values to probabilities, it cannot capture complex patterns and interactions present in non-linear data [5].
3. **Feature Engineering:** Non-linear data often requires feature transformations or the use of non-linear models (like SVMs or Neural Networks) to effectively learn and represent the underlying patterns [3].

Therefore, Logistic Regression is best suited for problems where the decision boundary can be approximated well by a linear function. For non-linear data, more complex models that can capture non-linear relationships are typically preferred.

## 13) Differentiate between Adaboost and Gradient Boosting.

To differentiate between Adaboost and Gradient Boosting:

1. **Base Learners:**
  - **AdaBoost:** Uses a series of weak learners (often decision trees) sequentially, each focusing on instances that were incorrectly predicted by the

previous learner. It adjusts the weights of incorrectly classified instances to emphasize the harder cases.

- **Gradient Boosting:** Also employs weak learners (usually decision trees) sequentially, but it fits each subsequent model to the residual errors of the previous one. It minimizes the error of the overall ensemble by using gradient descent-like techniques.

## 2. Training Process:

- **AdaBoost:** Adjusts the weights of instances at each step based on how well they were classified by previous models, aiming to improve performance on difficult instances.
- **Gradient Boosting:** Uses gradient descent techniques to minimize the overall error of the model by adding models sequentially, each correcting errors made by its predecessors.

## 3. Handling of Errors:

- **AdaBoost:** Focuses on instances that were misclassified by earlier models, giving them higher weights in subsequent iterations to improve performance on these instances.
- **Gradient Boosting:** Models are trained sequentially to minimize the residual errors (difference between predicted and actual values), effectively reducing the overall error of the ensemble.

## 4. Complexity and Robustness:

- **AdaBoost:** Can be sensitive to noisy data and outliers due to its iterative focusing on difficult instances, which can lead to overfitting in some cases.
- **Gradient Boosting:** Generally, more robust to noisy data and outliers because it optimizes the model by minimizing errors rather than focusing on instance weights.

# 14) What is bias-variance trade off in machine learning?

The bias-variance tradeoff in machine learning refers to a fundamental dilemma encountered when developing predictive models. It involves finding a balance between two sources of error:

## 1. Bias:

- Bias is the error introduced by approximating a real-world problem with a simplified model. It represents how closely the model's predictions match the true values.
- Models with high bias are too simplistic and may consistently miss relevant patterns in the training data. They typically underfit the data.

## 2. Variance:

- It measures the model's sensitivity to small fluctuations in the training data. It reflects how much the model's predictions vary for different training sets.

- Models with high variance are overly complex and tend to capture noise in the training data, leading to overfitting. They perform well on training data but poorly on unseen data.

**Tradeoff:**

- Increasing model complexity typically reduces bias but increases variance, and vice versa. The goal is to find the optimal balance where the total error (sum of bias and variance) is minimized, leading to better generalization on unseen data.

Understanding and managing the bias-variance tradeoff is crucial for selecting appropriate algorithms, tuning model parameters, and preventing overfitting or underfitting in machine learning applications.

**15) Give short description each of Linear, RBF, Polynomial kernels used in SVM.**

**1. Linear Kernel:**

- The Linear kernel computes the dot product between two vectors. It is suitable for linearly separable data or when the number of features is very large.
- Effective when the decision boundary is expected to be linear.

**2. RBF (Radial Basis Function) Kernel:**

- The RBF kernel computes similarity based on the Euclidean distance between data points in a high-dimensional space. It is capable of modeling complex, nonlinear decision boundaries.
- Widely used for classification and regression tasks where the data is not linearly separable.

**3. Polynomial Kernel:**

- The Polynomial kernel computes the similarity between vectors using a polynomial function. It allows for more flexibility in defining the decision boundary compared to the linear kernel.
- Useful for capturing higher-dimensional relationships in the data, though it requires careful tuning of parameters like degree and coefficient.

These kernels enable SVMs to handle data that may not be linearly separable in the original feature space by transforming it into a higher-dimensional space where separation is possible.